

Exercises

Import data

3.2.5: 1,4,5

1

```
# Combining all conditions into single pipe for fun
a <- flights |>
  filter(arr_delay>=2) |>
  filter(dest == 'IAH' | dest == 'HOU') |>
  filter(carrier == 'AA' | carrier == 'UA' | carrier == 'DL') |>
  filter(month == 7 | month == 8 | month == 9) |>
  filter(dep_delay<=0)

# Have to do this pipe separate cause it contradicts if in other one
b <- flights |>
  filter(dep_delay>=1 & arr_delay<30)
```

4

```
un <- flights |>
  distinct(month,day)

dim(un)
```

```
[1] 365  2
```

There are 365 unique month, day pairs, so yes there was a flight everyday of 2013.

5

```
long <- flights |>
  arrange(desc(distance)) |>
  head(1)
```

```
short <- flights |>
  arrange(distance) |>
  head(1)
```

long

```
# A tibble: 1 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
1  2013     1     1     857           900        -3    1516         1530
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

short

```
# A tibble: 1 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
1  2013     7    27      NA           106         NA      NA           245
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

JFK to HNL is the longest and JFK to LGA is the shortest

3.3.5: 1,4

1

```
flights |>
  mutate(delay = dep_time-sched_dep_time) |>
  mutate(diff = delay-dep_delay) |>
  filter(diff != 0 | !is.na(diff))
```

```
# A tibble: 328,521 x 21
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517             515           2     830           819
2  2013     1     1     533             529           4     850           830
3  2013     1     1     542             540           2     923           850
4  2013     1     1     544             545          -1    1004          1022
5  2013     1     1     554             600          -6     812           837
6  2013     1     1     554             558          -4     740           728
7  2013     1     1     555             600          -5     913           854
8  2013     1     1     557             600          -3     709           723
9  2013     1     1     557             600          -3     838           846
10 2013     1     1     558             600          -2     753           745
# i 328,511 more rows
# i 13 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>, delay <int>, diff <dbl>
```

I would expect `dep_delay = dep_time - sched_dep_time`, but we can see they are actually storing time as `hmm`. So 558 is 5:58 and 600 is 6:00, so the intuitive approach doesn't work.

4

`any_of()` selects any variable in a vector. It helps when you want variables removed because calling it twice doesn't error.

3.5.7: 1,2,4,6

1

```
flights |>
  group_by(carrier) |>
  summarise(avg_dep_delay = mean(dep_delay, na.rm = TRUE)) |>
  arrange(desc(avg_dep_delay))
```

```
# A tibble: 16 x 2
  carrier avg_dep_delay
  <chr>         <dbl>
1 F9           20.2
```

2	EV	20.0
3	YV	19.0
4	FL	18.7
5	WN	17.7
6	9E	16.7
7	B6	13.0
8	VX	12.9
9	OO	12.6
10	UA	12.1
11	MQ	10.6
12	DL	9.26
13	AA	8.59
14	AS	5.80
15	HA	4.90
16	US	3.78

F9 is the worst.

2

```
flights |>
  group_by(dest) |>
  slice_max(dep_delay, n = 1, with_ties = FALSE) |>
  ungroup()
```

A tibble: 105 x 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	12	14	2223	2001	142	133	2304
2	2013	7	23	1139	800	219	1250	909
3	2013	1	25	123	2000	323	229	2101
4	2013	8	17	1740	1625	75	2042	2003
5	2013	7	22	2257	759	898	121	1026
6	2013	7	10	2056	1505	351	2347	1758
7	2013	6	14	1158	816	222	1335	1007
8	2013	2	21	1728	1316	252	1839	1413
9	2013	12	1	1504	1056	248	1628	1230
10	2013	4	10	25	1900	325	136	2045

i 95 more rows

i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,

```
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

4

It excludes the n min/max rows and returns the rest.

6

```
df <- tibble(
  x = 1:5,
  y = c("a", "b", "a", "a", "b"),
  z = c("K", "K", "L", "L", "K")
)
```

group_by() will set the metadata group values for the tibble for the columns passed.

```
df |>
  group_by(y)
```

```
# A tibble: 5 x 3
# Groups:   y [2]
   x y     z
<int> <chr> <chr>
1     1 a     K
2     2 b     K
3     3 a     L
4     4 a     L
5     5 b     K
```

arrange(y) will order y alphabetically

```
df |>
  arrange(y)
```

```
# A tibble: 5 x 3
   x y     z
<int> <chr> <chr>
```

1	1	a	K
2	3	a	L
3	4	a	L
4	2	b	K
5	5	b	K

This will get the mean x value for each y value:

```
df |>
  group_by(y) |>
  summarize(mean_x = mean(x))
```

```
# A tibble: 2 x 2
  y      mean_x
  <chr>   <dbl>
1 a      2.67
2 b      3.5
```

This will do the same but for each unique y,z pair (and drop the group metadata):

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x), .groups = "drop")
```

```
# A tibble: 3 x 3
  y      z      mean_x
  <chr> <chr>   <dbl>
1 a     K      1
2 a     L     3.5
3 b     K     3.5
```

These will get the mean_x for each unique y,z pair. The second one adds a column to the original tibble and the first just lists all unique values.

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x))
```

`summarise()` has grouped output by 'y'. You can override using the `.groups` argument.

```
# A tibble: 3 x 3
# Groups:   y [2]
  y     z   mean_x
  <chr> <chr> <dbl>
1 a     K     1
2 a     L     3.5
3 b     K     3.5
```

```
df |>
  group_by(y, z) |>
  mutate(mean_x = mean(x))
```

```
# A tibble: 5 x 4
# Groups:   y, z [3]
  x y     z   mean_x
  <int> <chr> <chr> <dbl>
1     1 a     K     1
2     2 b     K     3.5
3     3 a     L     3.5
4     4 a     L     3.5
5     5 b     K     3.5
```